

Monitoring and Tuning Oracle RAC Database

By Ahmed Baraka

Objectives

In this lecture, you should understand and describe the following:

- Performance tuning tools
- Global Cache wait events
- Global Enqueue wait events
- Common RAC-related performance issues and their resolutions
- Tips to tune Oracle RAC performance
- Using ADDM in RAC



Ahmed Baraka
Oracle Database Administrator

Tuning Database Performance: Overview

- Performance Tuning areas:
 - Application and SQL Tuning
 - Operating System
 - Database Objects and configuration settings
 - Hardware
 - Network
- Lecture scope: RAC-specific tuning topics



Ahmed Baraka
Oracle Database Administrator

Performance Tuning in Oracle RAC

- Tuning like a single-instance first
- RAC-specific tuning areas:
 - Single instance bottlenecks get exacerbated in RAC
 - Interconnect traffic
- Database performance tuning tools:
 - Performance views v\$ (specially those that retrieve wait events and System, RAC, and instance statistics)
 - Statspack or AWR reports
 - Enterprise Manager (Cloud Control or Express)



Ahmed Baraka
Oracle Database Administrator

RAC-specific Wait Events

- RAC-related wait events come under the category “Cluster Wait Class”.
 - A significant rise in a wait event could be a symptom to lead you to a performance issue root cause
 - Some wait events and statistics pattern are symptoms of known performance issues in RAC.
 - Used by ADDM to enable cache fusion diagnostics



Ahmed Baraka
Oracle Database Administrator

Wait Event Views

View	Description
V\$SYSTEM_EVENT	Total wait for an event since last startup
V\$SESSION_EVENT	Waits by session
V\$ACTIVE_SESSION_HISTORY	Snapshots of active database sessions taken once a second
V\$SQL	Statistics on shared SQL areas, it contains one row for each SQL statement
V\$SQLAREA	Statistics on shared SQL areas, group by SQL text
V\$SQLSTATS	Basic performance statistics for SQL, one row per SQL statement



Ahmed Baraka
Oracle Database Administrator

Global Cache Wait Events

gc [current/cr] block [2/3]-way

- Received after two or three network hops with no wait

gc [current/cr] block busy

- Received but not sent immediately

gc [current/cr] grant 2-way

- Not received and not mastered locally. Grant received immediately. If 90% of this wait event is less than 2ms, network response time is fine.

gc current grant busy

- Not received and not mastered locally. Grant received with delay.



Ahmed Baraka
Oracle Database Administrator

Global Cache Wait Events (cont..)

gc [current/cr] [block/grant] congested

- Block or grant received with delay because of CPU or memory lack

gc [current/cr] [failure/retry]

- Not received because of failure

gc buffer busy

- Block arrival time less than buffer pin time

gc current/cr request

- A gc request for a cr block or current buffer is in progress. The waits act as placeholders until the request completes



Global Cache Wait Events (cont..)

`gc remaster - gcs drm server freezer - gc quiesce`

- Instance remastering is going on




Ahmed Baraka
Oracle Database Administrator

Global Enqueue Wait Events

- Enqueue wait is not RAC specific, but it involves a global lock operation in RAC. Following are common enqueue types:

Enqueue Type	Description
TX	Transaction enqueue: used to protect from simultaneous update operations
TM	Table or partition enqueue: used to protect table definitions during DML operations
HW	High-water mark: acquired to synchronize a new block operation
SQ	Sequence: used to serialize concurrent request to a sequence number increments
US	Undo segment: mainly used by the Automatic Undo Management
TA	Enqueue used mainly for transaction recovery during instance recovery



Ahmed Baraka
Oracle Database Administrator

Enqueue Statistics

- Enqueue events reported in all wait event views, AWR, and Statspack reports.
- **V\$ENQUEUE_STATISTICS** can be queried to determine which enqueue has the highest impact on database service times and, eventually, response times.



Ahmed Baraka
Oracle Database Administrator

Session and System Statistics

View	Description
V\$SYSSTAT	Retrieves system-wide statistics Can be used to categorize the system statistics
V\$SESSTAT	Retrieves user session statistics. Usually linked with V\$STATNAME to obtain the statistic name. Can be used to obtain the top sessions
V\$INSTANCE_CACHE_TRANSFER	Retrieves statistics for the cache blocks transferred among instances through the interconnect



Ahmed Baraka
Oracle Database Administrator

Common RAC Tuning Tips

Consider application side tuning:

- Reduce long full-table scans in OLTP systems
- Enable and rise sequence caches
- Use partitioning to distribute the workload on instances
- Avoid unnecessary parsing
- Monitor unselective indexes and study how to remove them
- Monitor the interconnect and study configuring it properly
- Minimize locking usage



Using Oracle Sequences in RAC

- High-insert applications lead to leaf block contention
- In RAC, when insert rate is high, this lead to a high rate of current and CR blocks transferred between nodes.
- Known solutions:
 - Enable and increase sequence caches

```
CREATE SEQUENCE seq_rac_test CACHE 1000;
```

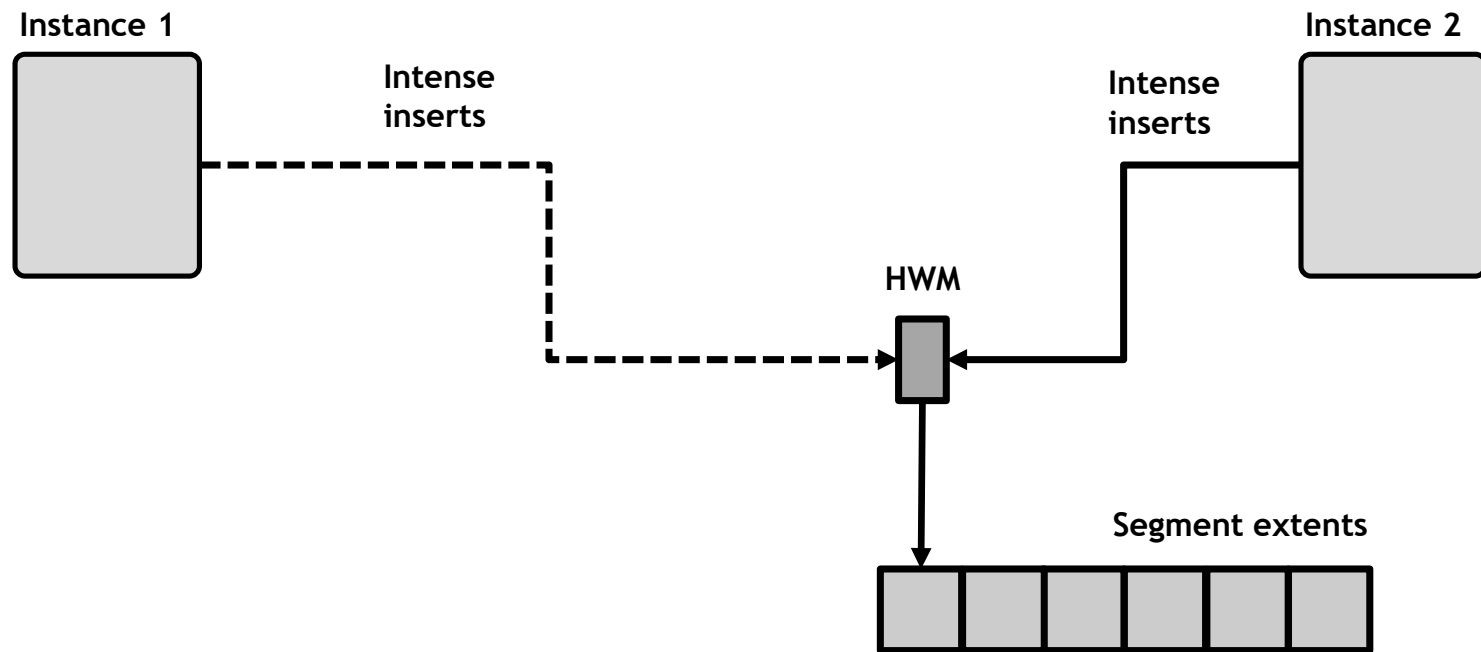
- Each instance use different range

```
CREATE TRIGGER insert_PKs ...  
  SELECT INSTANCE_NUMBER ... FROM V$INSTANCE
```




Ahmed Baraka
Oracle Database Administrator

High-watermark Enqueue



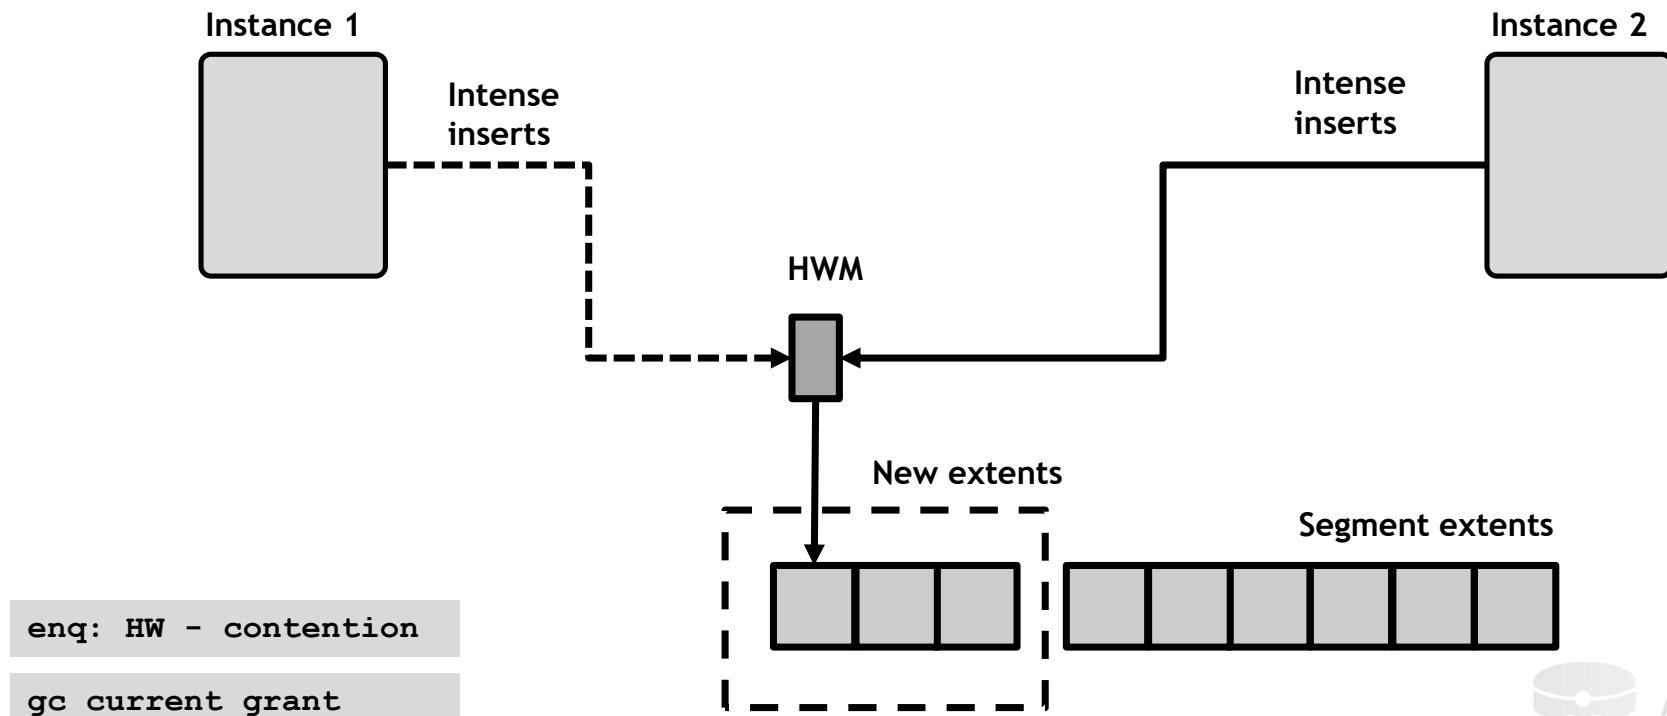
Ahmed Baraka
Oracle Database Administrator

High-watermark Enqueue



Ahmed Baraka
Oracle Database Administrator

High-watermark Enqueue



Ahmed Baraka
Oracle Database Administrator

About High-watermark Contention

- High-water mark contention characteristics:
 - Heavy inserts from RAC instances into the same segment at the HWM
 - New extents have to be created
 - Search for free space
- The common symptoms of the issue:
 - `enq: HW - contention`
 - `gc current grant` events
- To defuse the issue:
 - Uniform large extent sizes to the segments
 - Use partitions



Ahmed Baraka
Oracle Database Administrator

Automatic Workload Repository (AWR) in RAC

- Performance statistics are gathered and stored every hour
- In RAC, performance statistics from all active instances are taken
- Data is stored for each instance (but not for the entire cluster)
- Reports can be generated on the instance level or database level
- AWR report has Cluster related sections
- AWR reports can be used for tuning cases
- Requires separate license. If not obtained, use Statspack instead



About Active Session History (ASH)

- Provides statistics about current active sessions
- Useful for diagnosing performance issues
- Any session connected to the database and not idle is *active*
- It is based on **V\$ACTIVE_SESSION_HISTORY**
 - Snapshots of active database sessions taken every second
 - Most columns are the same as in **V\$SESSION** view
- ASH report can be obtained from OEM or generated in SQL*Plus:

```
$ORACLE_HOME/rdbms/admin/ashrpt.sql
```



Automatic Database Diagnostic Monitor (ADDM) for RAC

- Analyzes AWR snapshots for any possible performance issues
- Runs automatically after ARW snapshot is taken
- In RAC, it presents its data in different granularity:
 - The entire cluster (Database ADDM for RAC mode)
 - A specific instance (Local ADDM mode)
 - Subset of database instances (Partial ADDM mode)
- Can be invoked from OEM or through **DBMS_ADVISOR** and **DBMS_ADDM**



Ahmed Baraka
Oracle Database Administrator

Issues Analyzed by ADDM

- Instance congestion (top instances affecting the entire RAC)
- Object contention (buffer busy, top objects, and so on)
- Top consumers
- Lost blocks
- Interconnect devices issues including their throughput
- Latency problems in interconnect



Ahmed Baraka
Oracle Database Administrator

Summary

In this lecture, you should have learnt how to describe the following:

- Performance tuning tools
- Global Cache wait events
- Global Enqueue wait events
- Common RAC-related performance issues and their resolutions
- Tips to tune Oracle RAC performance
- Using ADDM in RAC



Ahmed Baraka
Oracle Database Administrator